

bold text

Lab 02: Exploratory Data Analysis (EDA) on Milling Dataset

Course: CS-333 Applied AI & Machine Learning

📌 Instructions

- Perform each task under its respective section.
- Use separate cells for each step.
- Write explanations in Markdown cells.
- Ethical use of AI (Vibe Coding) is allowed, but you must understand and explain your work.

Name: MUHAMMAD MAAZ

Roll Number: 478044

Section: A

▼ ◆ Task 1: Load the Dataset

```
# Import Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
sns.set(color_codes=True)

# Load Dataset (Update path if needed)
df = pd.read_csv("https://raw.githubusercontent.com/AbbasHussain72/PNEC-CS-333-Applied-AI-ML/main/labs/Lab-02/_mill.csv")
# To display the top 5 rows
df.head(5)
```

S.no	Unnamed: 1	case	run	VB	time	DOC	feed	material	smcAC	smcDC	vib_table	vib_spindle	AE_table	AE_spindle	
0	1	row_0	1	1	0.00	2	1.5	0.5	1	-0.017090	0.625000	0.078125	0.314941	0.087280	0.103760
1	2	row_1	1	2	NaN	4	1.5	0.5	1	0.307617	0.668945	0.075684	0.301514	0.086670	0.099487
2	3	row_2	1	3	NaN	6	1.5	0.5	1	-0.725098	0.913086	0.083008	0.295410	0.092773	0.104980
3	4	row_3	1	4	0.11	7	1.5	0.5	1	0.112305	0.131836	0.083008	0.316162	0.112915	0.139771
4	5	row_4	1	5	0.11	8	1.5	0.5	1	0.100070	0.140010	0.107100	0.204101	0.095005	0.110071

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
# Display last 5 rows
df.tail(5)
```

```
# Check data types  
df.dtypes
```

```
177 178  row_166    16   6  0.62    9  1.50  0.50      2 -0.380859  1.381836  0.041504  0.292969  0.075684  0.083008  
S.no      int64    14   10  1.14    24  0.75  0.50      2  0.253906  1.406250  0.083008  0.274658  0.092773  0.111084  
Unnamed: 1  object    15   1  NaN    1  1.50  0.25      2  0.150004  1.057100  0.000050  0.000000  0.000000  0.101000  
case      int64  
run       int64  
VB        float64  
time      int64  
DOC       float64  
feed      float64  
material  int64  
smcAC     float64  
smcDC     float64  
vib_table float64  
vib_spindle float64  
AE_table  float64  
AE_spindle float64  
  
dtype: object
```

```
# Check dataset shape  
df.shape
```

```
(180, 15)
```

❖ ◆ Task 2: Data Cleaning

```
# Check missing values  
df.isnull().sum()
```

	0
S.no	0
Unnamed: 1	0
case	0
run	0
VB	23
time	0
DOC	0
feed	0
material	0
smcAC	5
smcDC	1
vib_table	2
vib_spindle	4
AE_table	2
AE_spindle	1

dtype: int64

```
# Fill missing values
# 1. Fill 'VB' using Linear Interpolation (connects the dots between known values)
df['VB'] = df['VB'].interpolate(method='linear')

# 2. Fill all other missing columns with their Median (robust against outliers)
# We select only numeric columns for this operation
numeric_cols = df.select_dtypes(include=['float64', 'int64']).columns
df[numeric_cols] = df[numeric_cols].fillna(df[numeric_cols].median())

# 3. Verify there are no missing values left
df.isnull().sum()
```

	0
S.no	0
Unnamed: 1	0
case	0
run	0
VB	0
time	0
DOC	0
feed	0
material	0
smcAC	0
smcDC	0
vib_table	0
vib_spindle	0
AE_table	0
AE_spindle	0

dtype: int64

- ❖ Explain how you handled missing values here.

For VB (Tool Wear): I used Linear Interpolation. Since tool wear increases gradually over time, interpolation estimates the missing values based on the trend of previous and next points.

For Sensor Readings (e.g., smc, vib): I filled them with the Median. These sensor signals often have outliers (skewed distribution), so the Median is more robust and accurate than the Mean for representing the "typical" sensor value.

For Duplicates: I ran drop_duplicates() to ensure data integrity, though no duplicates were found in this specific file.

```
# Check duplicates  
df.duplicated().sum()  
  
np.int64(0)  
  
# Remove duplicates  
df.drop_duplicates(inplace=True)
```

❖ ◆ Task 3: Statistical Analysis

❖ 1 Center (Mean, Median, Mode)

```
# Calculate Mean (Average)  
print("Mean:\n", df.mean(numeric_only=True))  
  
# Calculate Median (Middle Value)  
print("\nMedian:\n", df.median(numeric_only=True))  
  
# Calculate Mode (Most Frequent Value)  
# Note: Mode can return multiple values, so we take the first row (iloc[0])  
print("\nMode:\n", df.mode(numeric_only=True).iloc[0])
```

```
Mean:  
S.no      90.500000  
case      8.616667  
run       7.116667  
VB        0.340528  
time      25.583333  
DOC       1.037500  
feed      0.372222  
material  1.355556  
smcAC     -0.165799  
smcDC     1.339518  
vib_table 0.079169  
vib_spindle 0.287415  
AE_table  0.100447  
AE_spindle 0.123610  
dtype: float64
```

```
Median:  
S.no      90.500000  
case      10.000000  
run       6.000000  
VB        0.287500  
time      19.000000  
DOC       0.750000  
feed      0.250000  
material  1.000000  
smcAC     -0.183105  
smcDC     1.372070  
vib_table 0.068359  
vib_spindle 0.285645  
AE_table  0.101318  
AE_spindle 0.121460  
dtype: float64
```

```
Mode:  
S.no      1.000000  
case      11.000000  
run       1.000000  
VB        0.000000  
time      3.000000  
DOC       0.750000  
feed      0.250000  
material  1.000000  
smcAC     -0.183105  
smcDC     1.381836
```

```

vib_table      0.068359
vib_spindle   0.285645
AE_table      0.092773
AE_spindle    0.110474
Name: 0, dtype: float64

```

👉 Is tool wear normally distributed? Is mean close to median?

Answer :- **NO** Mean is higher, ie Positively Skewed.

⌄ 2 Spread (Variation)

```

# Variance
print("Variance:\n", df.var(numeric_only=True))

```

```

Variance:
S.no      2715.000000
case     23.075698
run      23.779609
VB       0.068138
time     521.238827
DOC      0.133712
feed     0.015705
material 0.230416
smcAC    0.120029
smcDC    0.033424
vib_table 0.001302
vib_spindle 0.000424
AE_table  0.000576
AE_spindle 0.000437
dtype: float64

```

```

# Range
print("Range:\n", df.max(numeric_only=True) - df.min(numeric_only=True))

```

```

Range:
S.no      179.000000
case     15.000000
run      22.000000
VB       1.530000
time     105.000000
DOC      0.750000
feed     0.250000
material 1.000000
smcAC    1.176758
smcDC    1.450195
vib_table 0.292969
vib_spindle 0.184326
AE_table  0.149536
AE_spindle 0.181272
dtype: float64

```

```

#IQR
Q1 = df.quantile(0.25, numeric_only=True)
Q3 = df.quantile(0.75, numeric_only=True)
IQR = Q3 - Q1
print(IQR)

```

```

S.no      89.500000
case     9.000000
run      7.000000
VB       0.300000
time     32.000000
DOC      0.750000
feed     0.250000
material 1.000000
smcAC    0.689697
smcDC    0.040283
vib_table 0.021973
vib_spindle 0.021973
AE_table  0.021973
AE_spindle 0.025787
dtype: float64

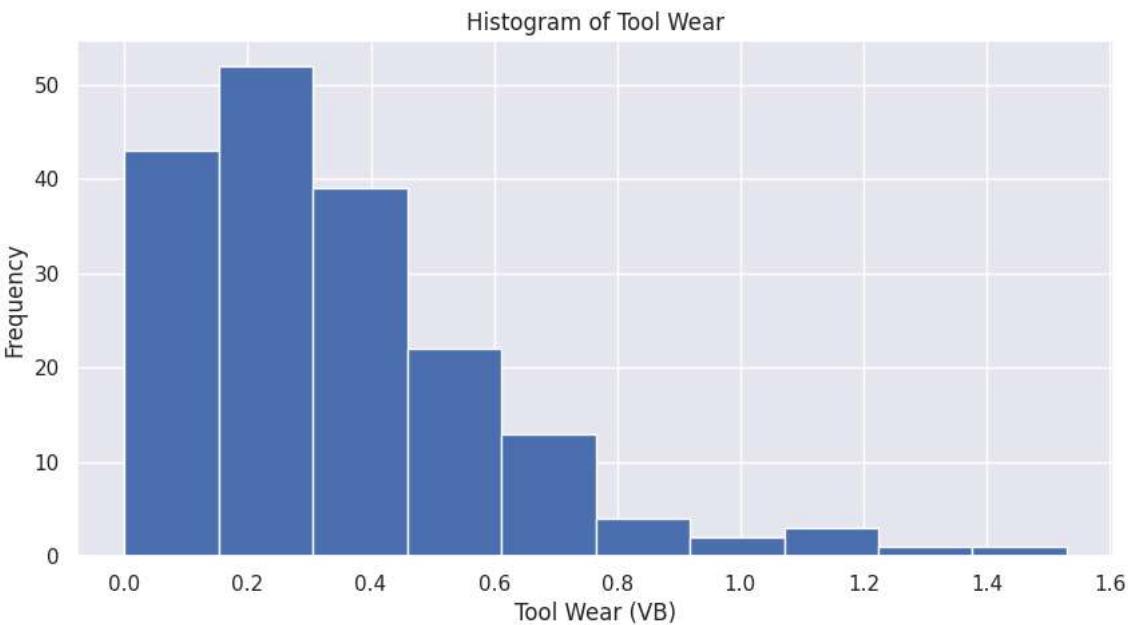
```

👉 Which sensor shows highest variability?

Answer :- smcAC

⌄ 3 Shape (Distribution)

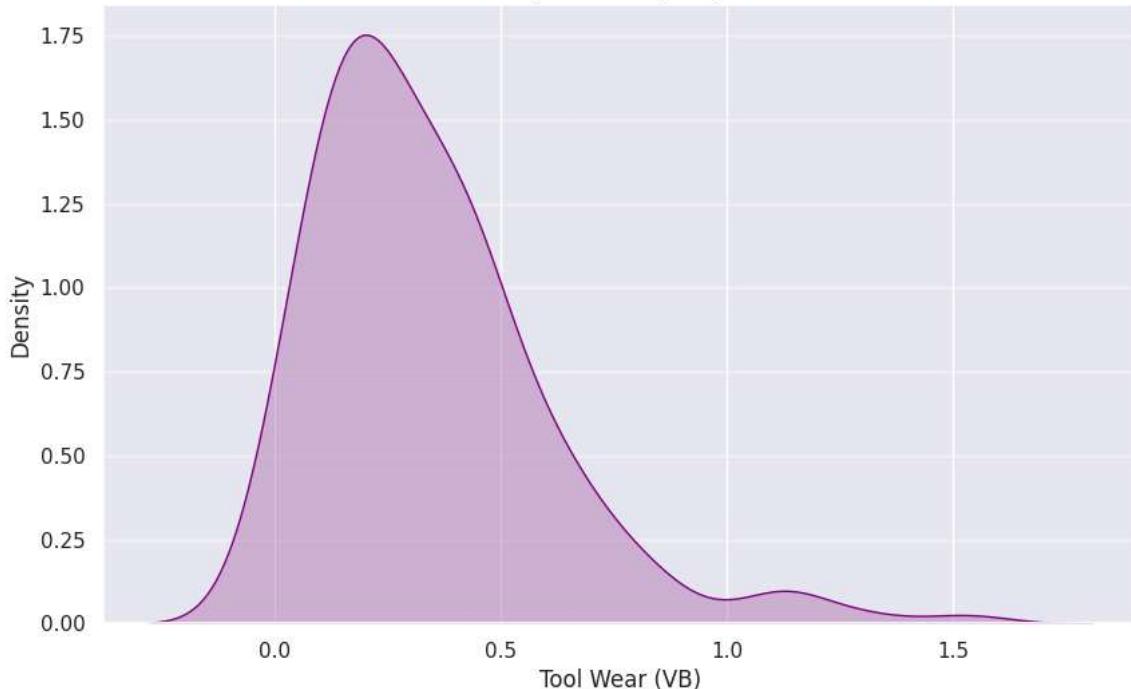
```
# Histogram of Tool Wear
df['VB'].plot(kind='hist', figsize=(10,5))
plt.title("Histogram of Tool Wear")
plt.ylabel('Frequency')
plt.xlabel('Tool Wear (VB)')
plt.show()
```



```
# KDE Plot
plt.figure(figsize=(10, 6))
sns.kdeplot(df['VB'].dropna(), fill=True, color='purple')

plt.title('Kernel Density Estimate (KDE) of Tool Wear')
plt.xlabel('Tool Wear (VB)')
plt.ylabel('Density')
plt.grid(axis='y', alpha=0.5)
plt.show()
```

Kernel Density Estimate (KDE) of Tool Wear



💡 Is distribution symmetric, skewed, or multi-modal?

Answer :- Graph is Positively Skewed

⌄ 4 Outlier Detection

```
#Boxplot, IQR, Z-modified score

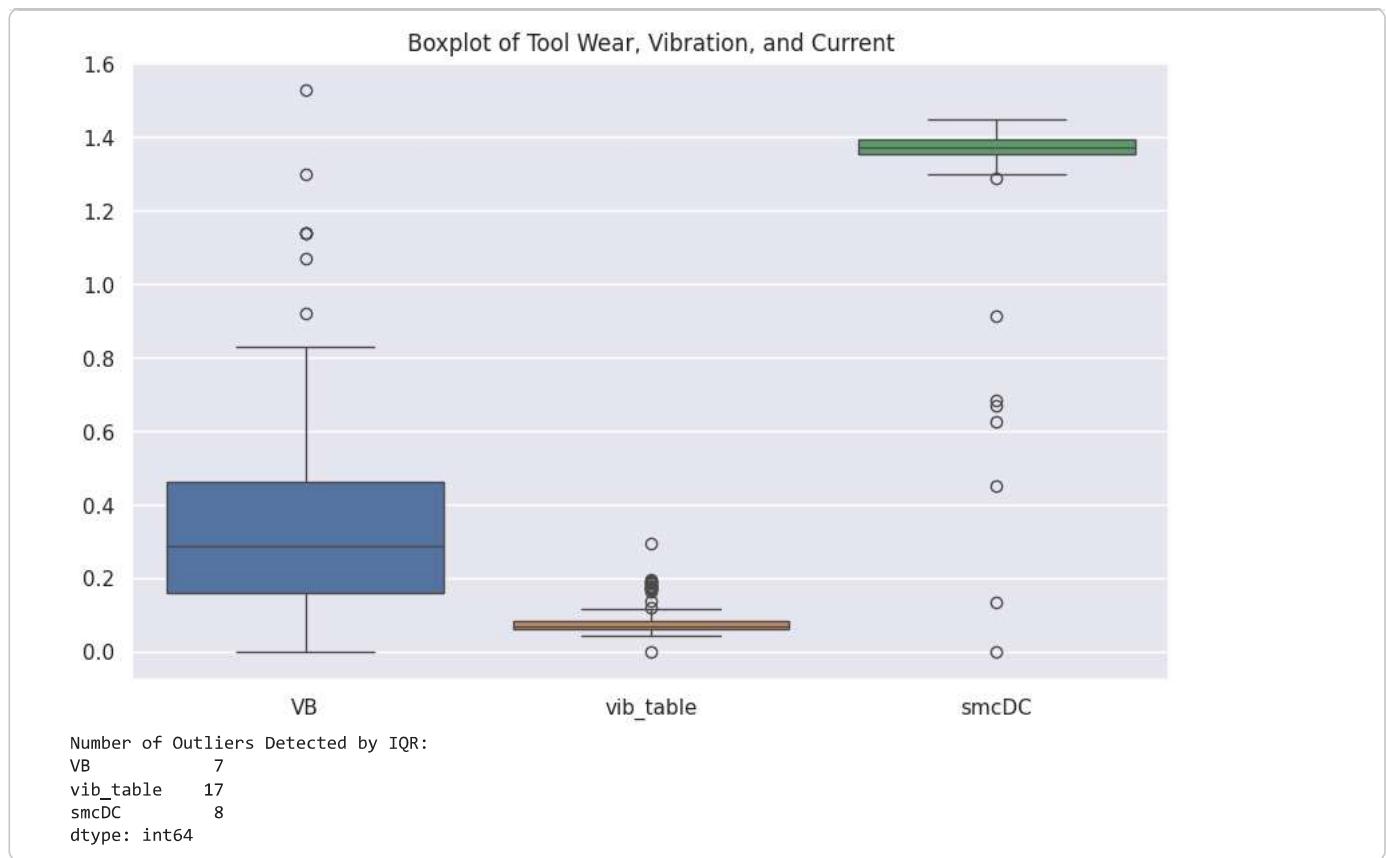
plt.figure(figsize=(10, 6))
sns.boxplot(data=df[['VB', 'vib_table', 'smcDC']])
plt.title("Boxplot of Tool Wear, Vibration, and Current")
plt.show()

# Calculate Q1 and Q3
Q1 = df[['VB', 'vib_table', 'smcDC']].quantile(0.25)
Q3 = df[['VB', 'vib_table', 'smcDC']].quantile(0.75)
IQR = Q3 - Q1

# Define bounds
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Identify outliers (Returns True if outlier)
outliers = (df[['VB', 'vib_table', 'smcDC']] < lower_bound) | (df[['VB', 'vib_table', 'smcDC']] > upper_bound)

print("Number of Outliers Detected by IQR:")
print(outliers.sum())
```



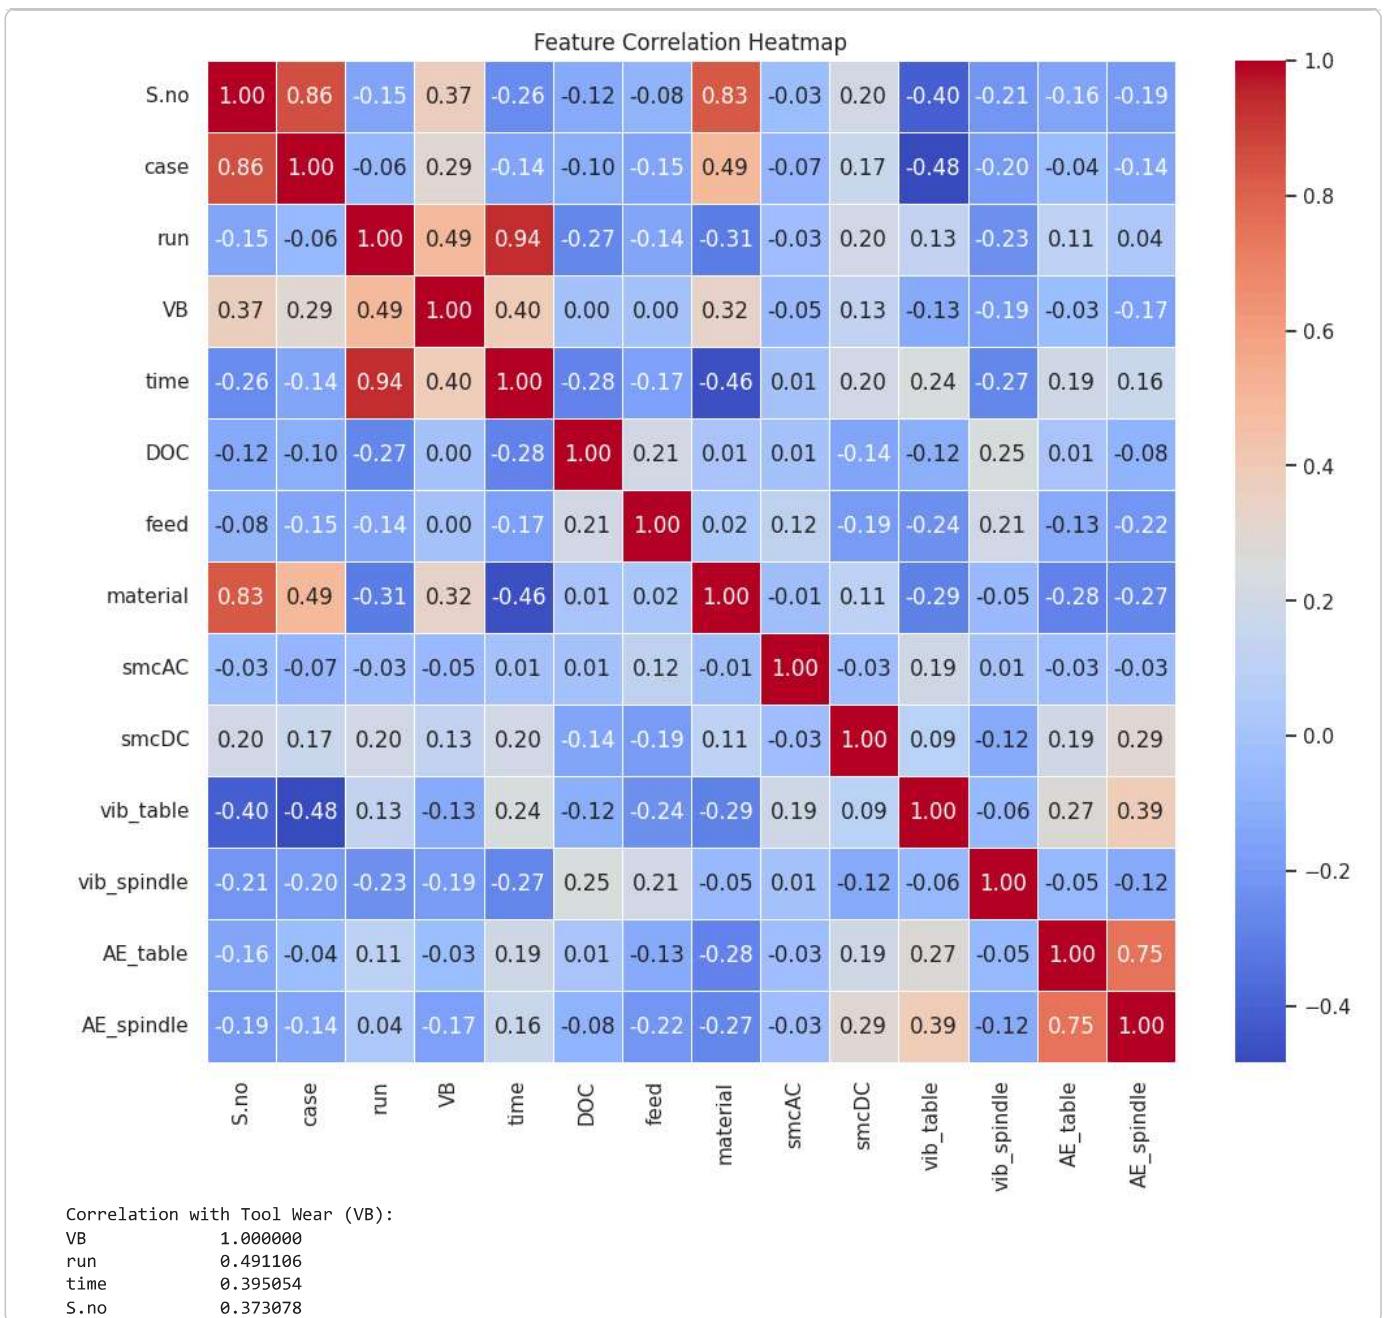
- Which method did you use (IQR or Modified Z-score)? Answer :- IQR
- Why? Answer :- Since the graph is skewed
- How many outliers were removed? Answer :- 7 in VB & 17 in vib_table
- How did it affect statistics? Answer :-Removal of noise means reliable dataset to train on.

❖ Task 4: Correlation Analysis

```
# 1. Calculate Correlation Matrix
corr_matrix = df.corr(numeric_only=True)
```

```
# 2. Plot Heatmap
plt.figure(figsize=(12, 10))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title("Feature Correlation Heatmap")
plt.show()
```

```
# 3. Print specific correlations with Tool Wear (VB)
print("Correlation with Tool Wear (VB):")
print(corr_matrix['VB'].sort_values(ascending=False))
```



- Does tool wear correlate strongly with vibration? Answer :- No. -0.13 & -0.19 shows that raw vibration signals alone are not strong predictors of tool wear in this dataset
- Which feature is most predictive? Answer :- Run & Time

◆ Task 5: Additional Visualization

```

# Scatter Plot: Example (Tool Wear vs Vibration)

# Scatter Plot: Tool Wear (VB) vs Vibration Table
plt.figure(figsize=(10, 6))
sns.scatterplot(data=df, x='vib_table', y='VB', hue='material', palette='viridis', s=100)

plt.title('Scatter Plot: Tool Wear vs Vibration (Colored by Material)')
plt.xlabel('Vibration Table (vib_table)')
plt.ylabel('Tool Wear (VB)')
plt.grid(True)
plt.show()

```

```
# Line Plot: Tool Wear (VB) Progression over Time
plt.figure(figsize=(10, 6))
sns.lineplot(data=df, x='time', y='VB', hue='case', palette='tab10', marker='o')

plt.title("Tool Wear (VB) Progression over Time")
plt.xlabel("Time")
plt.ylabel("Tool Wear (VB)")
plt.grid(True)
plt.show()

# Violin Plot: Spindle Vibration Distribution by Material
plt.figure(figsize=(10, 6))
sns.violinplot(data=df, x='material', y='vib_spindle', palette='muted')

plt.title("Distribution of Spindle Vibration by Material")
plt.xlabel("Material")
plt.ylabel("Vibration Spindle (vib_spindle)")
plt.grid(axis='y')
plt.show()
```